



Hands-on TensorFlow 2.0

Josh Gordon (@random_forests)

Amit Patankar (@av8ramit24)

At SciPy Tokyo

Slides from today
bit.ly/scipy-slides





Hands-on workshop

You will need

- A laptop with an internet connection.
- There's nothing to install in advance.

Agenda

- Beginner exercises in the first half
- Advanced examples in the second

Deep Learning is a huge space (our goal is not to cover everything, just to get you started).



TensorFlow

An open source Deep Learning library

- Released by Google in 2015
- **>1800** contributors worldwide

TensorFlow 2.0 (we'll use this today!)

- **Easier to use**
- Code styles for beginners and experts
- Alpha released in March, 2019



The person is riding a surfboard in the waves.



Exercises

Exercises

- Installing TF2 and using Colab
- Linear regression
- MNIST (with Keras Sequential)
- MNIST (with Keras Subclassing)
- Structured data

Advanced

- Deep Dream
- Neural machine translation
- Image Colorization



Topics

For beginners and experts

- Keras Sequential
- Keras Subclassing
- Built-in vs custom training loops

Beyond Hello World

- Interlingual representations

Under the hood

- AutoGraph and tf.function
- TF2 vs TF1

Learning more

- Book recommendations

Exercise 1

Linear regression in TensorFlow 2.0



Exercise 1

Goals

- Install TensorFlow 2.0
- Introduce Colab
- **Introduce ingredients** (predict, loss, improve, repeat)

Visit

bit.ly/tf-ws1



Why is Python popular for scientific computing?

—



Ballpark benchmarks

About how much slower is Python than C?



Ballpark benchmarks

About how much slower is Python than C?

- Multiplying matrices: +/- 100X
- 6 seconds vs. 10 minutes
- Running vs. flying (6 MPH and 600 MPH)

Python is a great choice for scientific computing

- Why?



Ballpark benchmarks

About how much slower is Python than C?

- Multiplying matrices: +/- 100X
- 6 seconds vs. 10 minutes
- Running vs. flying (6 MPH and 600 MPH)

Python is a great choice for scientific computing

- Why?

NumPy

- C performance, Python ease of use



TensorFlow is basically

NumPy

- + GPU / TPU support
- + AutoDiff
- + Utilities to help you write neural networks (layers, optimizers)

TensorFlow

- A C++ engine to accelerate code written in Python.
- **Bonus:** compiled to a graph that can run on devices **without a Python interpreter** (phones, web browsers)

You can use TF 2.0 like NumPy

```
import tensorflow as tf # Assuming TF 2.0 is installed

a = tf.constant([[1, 2],[3, 4]])
b = tf.matmul(a, a)

print(b)
# tf.Tensor( [[ 7 10] [15 22]], shape=(2, 2), dtype=int32)

print(type(b.numpy()))
# <class 'numpy.ndarray'>
```



For beginners and experts

For beginners

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

TF 1.x

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```

TF 2.0

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(512, activation='relu'),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Dense(10, activation='softmax')
])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

model.fit(x_train, y_train, epochs=5)
model.evaluate(x_test, y_test)
```





Keras and tf.keras

In my view, the **clearest Deep Learning library** that exists today.

- For fast prototyping, advanced research, and production.

keras.io = reference implementation

- `import keras`

tf.keras = TensorFlow's implementation (a superset, built-in to TF, no need to install Keras separately)

- `from tensorflow import keras`

Exercise 2

Fashion MNIST in TensorFlow 2.0



Exercise 2

Goals

- Learn about the Sequential API
- Train a simple image classifier

Visit

bit.ly/tf-ws4



playground.tensorflow.org

For experts

```
class MyModel(tf.keras.Model):  
    def __init__(self, num_classes=10):  
        super(MyModel, self).__init__(name='my_model')  
        self.dense_1 = layers.Dense(32, activation='relu')  
        self.dense_2 = layers.Dense(num_classes, activation='sigmoid')  
  
    def call(self, inputs):  
        # Define your forward pass here,  
        x = self.dense_1(inputs)  
        return self.dense_2(x)
```



What's the difference?



Symbolic vs Imperative APIs

Symbolic (Keras Sequential)

- Your model is a graph of layers
- Any graph you compile will run
- **TensorFlow helps you debug** by catching errors at **compile time**



Symbolic vs Imperative APIs

Symbolic (Keras Sequential)

- Your model is a graph of layers
- Any graph you compile will run
- **TensorFlow helps you debug** by catching errors at **compile time**

Imperative (Keras Subclassing)

- Your model is Python bytecode
- Complete flexibility and control
- Harder to debug / **harder to maintain**

Use a built-in training loop...

```
model.fit(x_train, y_train, epochs=5)
```

Or define your own

```
model = MyModel()
```

```
with tf.GradientTape() as tape:  
    logits = model(images)  
    loss_value = loss(logits, labels)
```

```
grads = tape.gradient(loss_value, model.trainable_variables)  
optimizer.apply_gradients(zip(grads, model.trainable_variables))
```

TensorBoard

```
tb_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir)
```

```
model.fit(  
    x_train, y_train, epochs=5,  
    validation_data=[x_test, y_test],  
    callbacks=[tb_callback])
```

- ☐ Show data download links
- ☒ Ignore outliers in chart scaling

Tooltip sorting method: **default**

Smoothing

0.6

Horizontal Axis

STEP

RELATIVE

WALL

Runs


Write a regex to filter runs

- ☒ 20190227-033014/test
- ☒ 20190227-033014/train

Filter tags (regular expressions supported)


accuracy

accuracy
tag: accuracy




loss

loss
tag: loss



sequential



Exercise 3

Keras model subclassing and TensorBoard



Exercise 3

Goals

- Learn about the Subclassing API
- See TensorBoard running in the browser

Visit

bit.ly/tf-ws3

Note: you may need to replace `tf-nightly` with `!pip install -q tensorflow-gpu==2.0.0-alpha0`



How to fix the TensorBoard example

```
!pip install -q tf-nightly-2.0-preview  
!pip install -q tensorboard==1.13.0  
%load_ext tensorboard.notebook
```

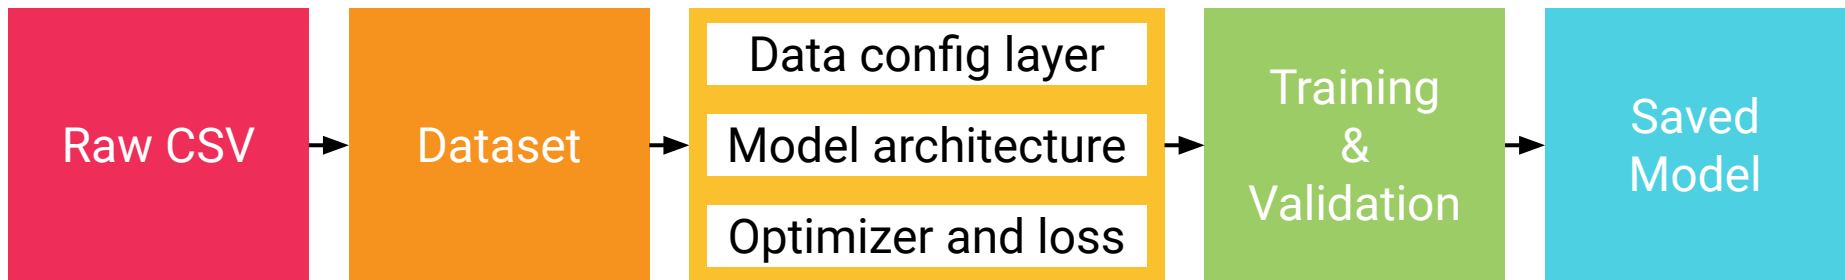
`bit.ly/tf-ws3`

(Thanks Amit!)



Structured data

Multi-stage Process






```
interestingNumbers = {  
    "Prime": [2, 3, 5, 7, 11, 13],  
    "Fibonacci": [1, 1, 2, 3, 5, 8],  
    "Square": [1, 4, 9, 16, 25],  
}
```

```
largest = 0  
for (kind, numbers) in interestingNumbers.items:  
    for x in numbers:  
        if x > largest:  
            largest = x  
  
print(largest)
```



Image of retina



Blood pressure predictions
focus on blood vessels



Our Dataset

Data: Heart Disease [V.A. Medical Center](#)

Task: Binary Classification (Healthy/Heart Disease)

Number of examples: ~300


Features:

- **Real:** *Age, Blood Pressure, Cholesterol*
- **Categorical - Int:** *Gender, EKG Results*
- **Categorical - String:** *Thallium heart scan*




Our Features Explained

1. **cp**: Chest pain type
2. **trestbps**: Resting blood pressure
3. **chol**: Serum cholesterol
4. **fbs**: Blood sugar > 120
5. **restecg**: Type of EKG result
6. **thalach**: Max heart rate achieved
7. **exang**: Exercise induced angina
8. **oldpeak**: ST depression (exercise induced)
9. **slope**: Slope of peak ST segment
10. **ca**: # of vessels colored by fluoroscopy
11. **thal**: Thallium heart scan results



The Data

age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
63	1	1	145	233	1	2	150	0	2.3	3	0	fixed	0
67	1	4	160	286	0	2	108	1	1.5	2	3	normal	1
67	1	4	120	229	0	2	129	1	2.6	2	2	reversible	0
37	1	3	130	250	0	0	187	0	3.5	3	0	normal	0
41	0	2	130	204	0	2	172	0	1.4	1	0	normal	0
56	1	2	120	236	0	0	178	0	0.8	1	0	normal	0
62	0	4	140	268	0	2	160	0	3.6	3	2	normal	1




Raw
CSV

Loading data

```
# Match CSV data types by column
defaults = [tf.int32, tf.int32 ... tf.string, tf.float32]


dataset = tf.contrib.data.CsvDataset(
    ['heart.csv.train'], defaults, header=True)
```



Loading data


```
print(list(dataset.take(1)))
```

```
[(<tf.Tensor: id=188, shape=(), dtype=int32, numpy=63>,  
  <tf.Tensor: id=189, shape=(), dtype=int32, numpy=1>,  
  <tf.Tensor: id=190, shape=(), dtype=int32, numpy=145>, ...  
  <tf.Tensor: id=191, shape=(), dtype=string, numpy='fixed'>,  
  <tf.Tensor: id=192, shape=(), dtype=float32, numpy=1>)]
```



Parsing data


```
def _parse_csv_row(*vals):  
    # Format each row for the model input  
  
    return features, labels
```



Parsing data

```
col_names = ['age', 'sex', 'cp', 'trestbps'...]
```


```
def _parse_csv_row(*vals):  
    # Format each row for the model input  
    # Element of val is a tensor  
    features = dict(zip(col_names, vals[:-1]))  
    labels = vals[-1]  
  
    return features, labels
```



Parsing data

```
dataset = dataset.shuffle(TRAINING_SIZE)
```


```
dataset = dataset.map(_parse_csv_row).batch(BATCH_SIZE)
```



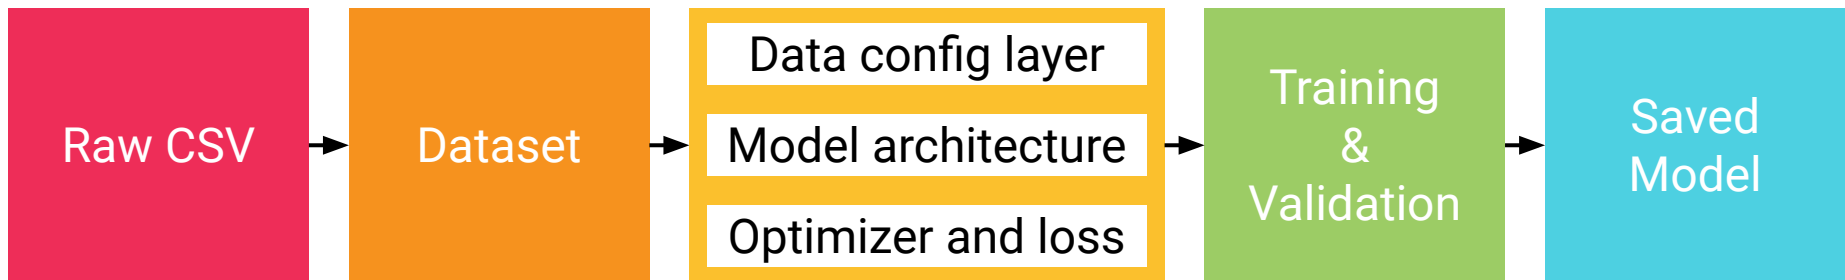
Parsing data

```
dataset = dataset.shuffle(TRAINING_SIZE)
dataset = dataset.map(_parse_csv_row).batch(BATCH_SIZE)
print(list(dataset.take(1)))
```


```
({'age': <tf.Tensor: shape=(64,), dtype=int32,
    array([47, ... 77, 32, 56])>,
  ...
  'thal': <tf.Tensor: shape=(64,), dtype=string
    array([[ 'reversible', ... 'normal' ]])>},
<tf.Tensor: shape=(64,), dtype=float64,
    array([0, 0, 1, ... 1, 0, 1])>)
```



Multi-stage Process




Multi-stage Process




Feature columns

- **Numeric columns**
 - Age, income, weight
- **Bucketized columns**
 - Decades, Age in ranges
- **Categorical identity columns**
 - Gender (0/1)
- **Categorical vocabulary column**
 - Countries (USA, Canada, Mexico)
- **Hashed column**
 - Object names
- **Crossed column**
 - Age along with gender




Defining categorical features

```
# thal / string / 1 of 3 values  
vocab = ['normal', 'fixed', 'reversible']  
thal_cc = tf.keras.feature_column.  
    categorical_column_with_vocabulary_list(  
        'thal', vocabulary_list=vocab)
```




Defining categorical features

```
# thal / string / 1 of 3 values
vocab = ['normal', 'fixed', 'reversible']
thal_cc = tf.keras.feature_column.
    categorical_column_with_vocabulary_list(
        'thal', vocabulary_list=vocab)
thal_embedding = tf.keras.feature_column.
    embedding_column(thal_cc, dimension=3)
```




Defining numeric features

```
# age / real integers  
age_nc = tf.keras.feature_column.  
    numeric_column('age')
```



Defining feature columns


```
feature_columns = [age_nc, sex_embedding, ...  
                  , thal_embedding]
```



Defining feature columns


```
feature_columns = [age_nc, sex_embedding, ...  
                  , thal_embedding]
```

```
feature_layer = tf.python.feature_column.  
    FeatureLayer(feature_columns)
```




Building a model

```
model = tf.keras.Sequential([  
    feature_layer,  
    tf.keras.layers.Dense(128, activation=tf.nn.relu),  
    tf.keras.layers.Dense(64, activation=tf.nn.relu),  
    tf.keras.layers.Dense(1, activation=tf.nn.sigmoid)  
])
```

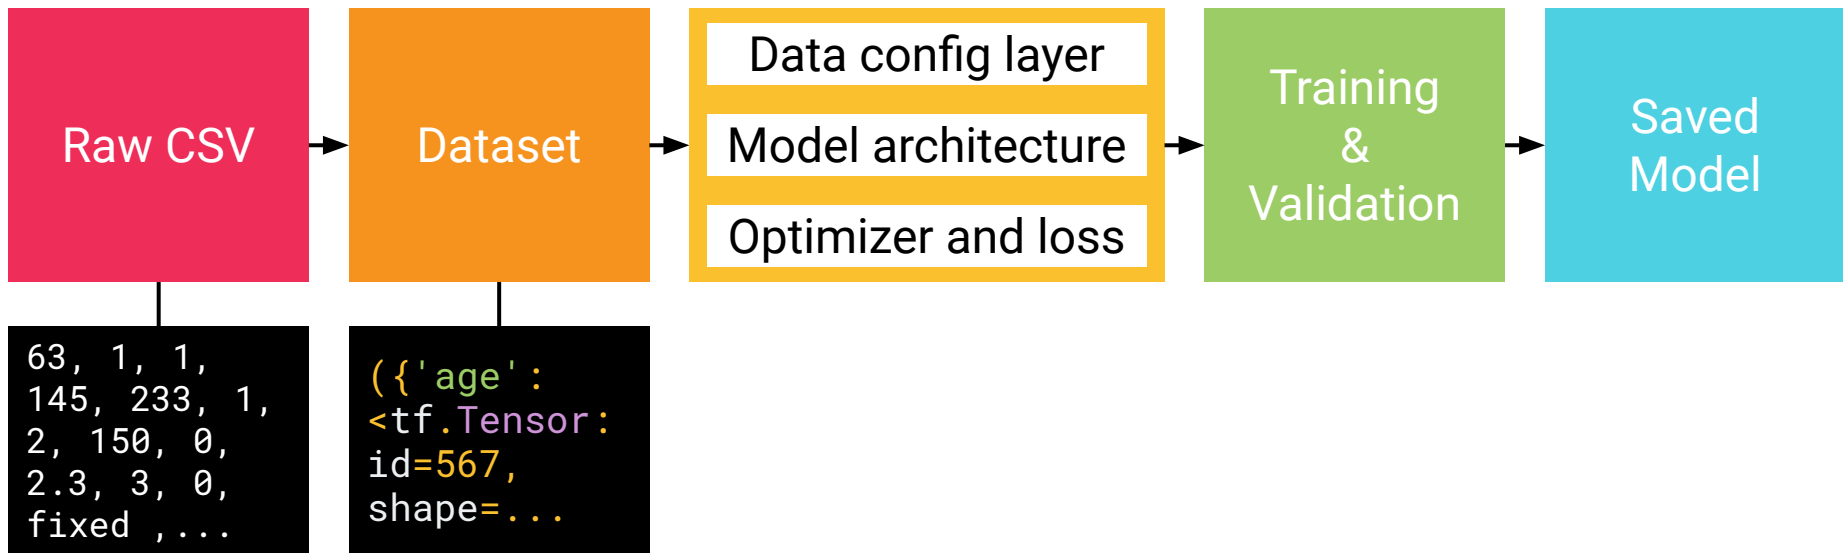


Building a model


```
model.compile(  
    optimizer=tf.train.AdamOptimizer(),  
    loss=tf.keras.losses.binary_crossentropy,  
    metrics=['accuracy'])
```



Multi-stage Process



Multi-stage Process




Training a model

```
for epoch in range(1,21):  
    # Print epoch, check validation metrics  
    model.fit(dataset,  
               steps_per_epoch=TRAINING_SIZE/BATCH_SIZE)
```


Epoch 20:

8/8 [=====] - 1s 41ms/step -
loss: 0.3475 - acc: 0.8433



Validating our model


```
test_ds = tf.contrib.data.CsvDataset(  
    ['heart.csv.test'], defaults)  
test_ds = test_ds.map(_parse_csv_row).batch(50)  
  
loss, accuracy = model.evaluate(test_ds, steps=1)  
print("Loss: {}\nAccuracy: {}".format(loss,  
accuracy))
```



Validating our model

```
test_ds = tf.contrib.data.CsvDataset(  
    ['heart.csv.test'], defaults)  
test_ds = test_ds.map(_parse_csv_row).batch(50)
```

```
loss, accuracy = model.evaluate(test_ds, steps=1)  
print("Loss: {}\nAccuracy: {}".format(loss,  
accuracy))
```




Validating our model

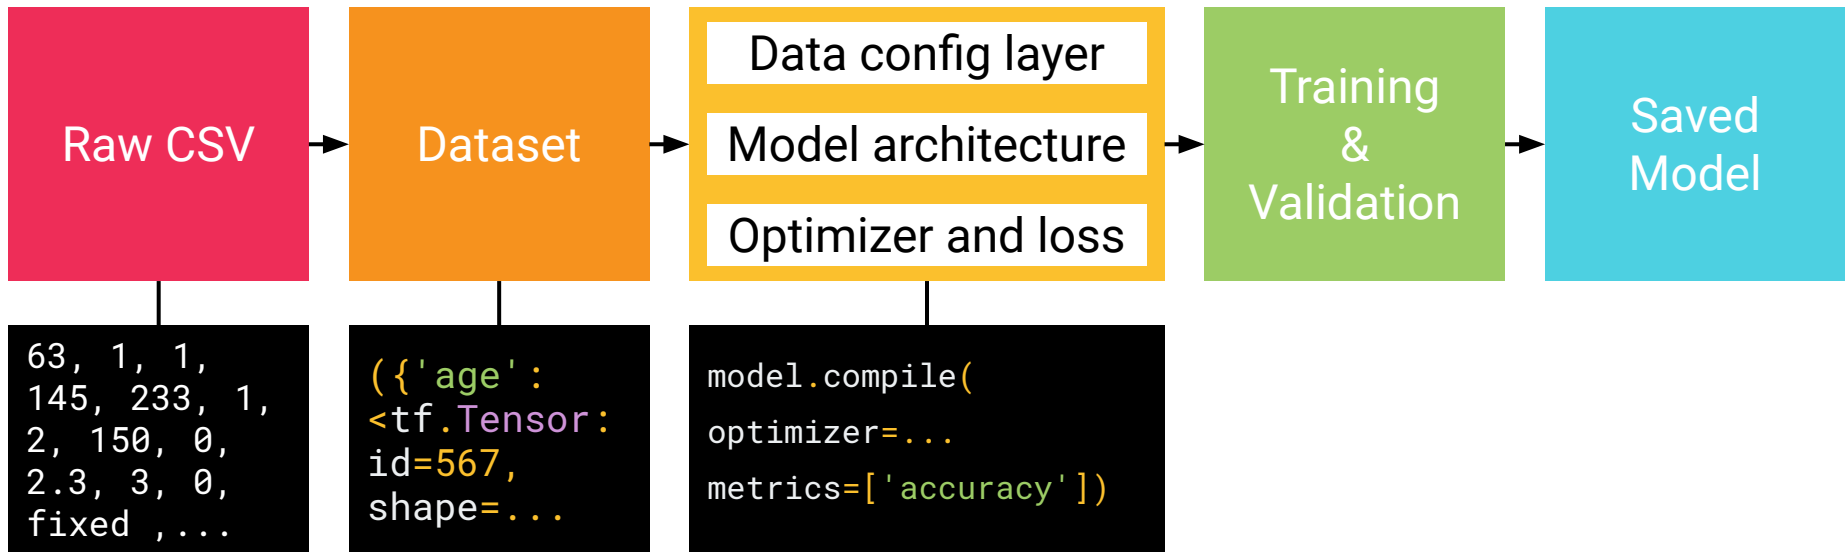
```
test_ds = tf.contrib.data.CsvDataset(  
    ['heart.csv.test'], defaults)  
test_ds = test_ds.map(_parse_csv_row).batch(50)  
  
loss, accuracy = model.evaluate(test_ds, steps=1)  
print("Loss: {}\nAccuracy: {}".format(loss,  
    accuracy))
```

Loss: 0.4622


Accuracy: 0.8400



Multi-stage Process




Multi-stage Process



Export to SavedModel

```
export_dir = tf.contrib.saved_model.  
    save_keras_model(model, 'keras_nn')
```


```
keras_nn/  
1536162174/  
    saved_model  
    variables/  
    assets/
```



Saved
Model

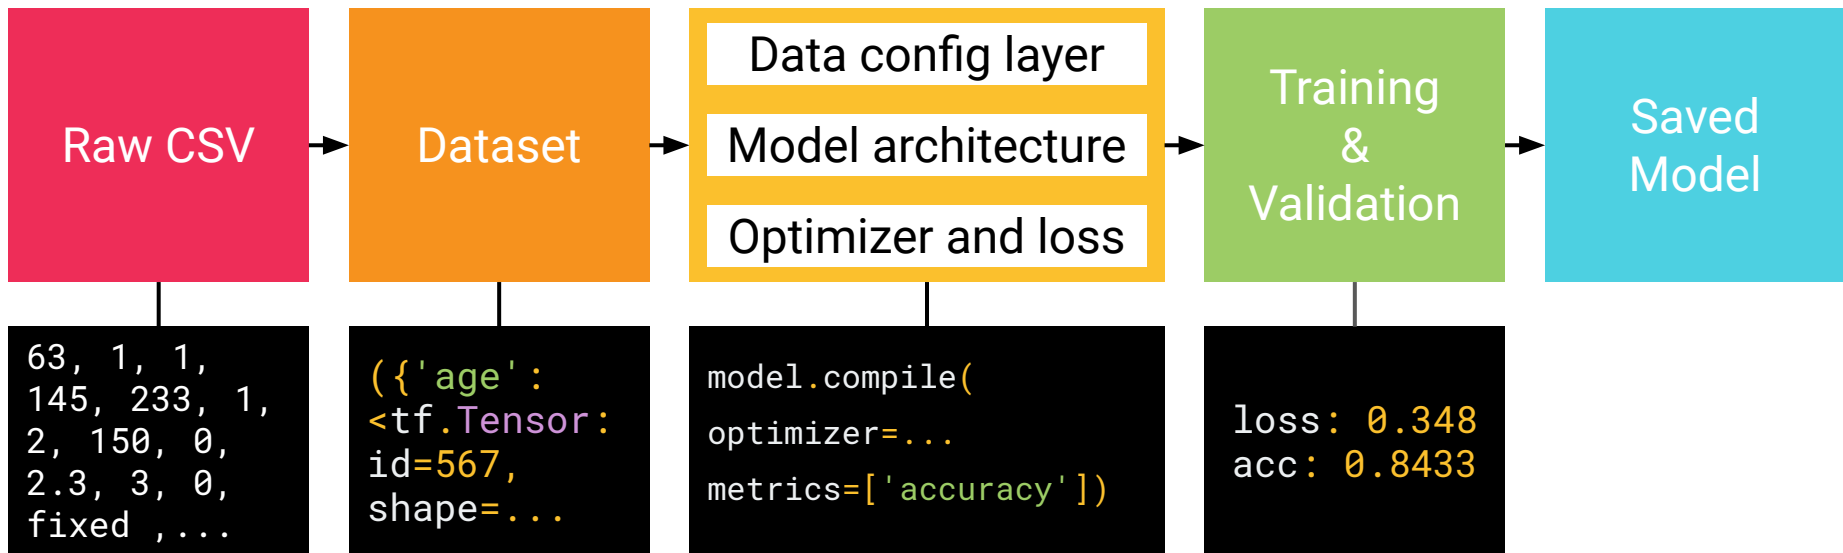
Restore from SavedModel

```
restored_model = tf.contrib.saved_model.  
    load_keras_model(export_dir)
```

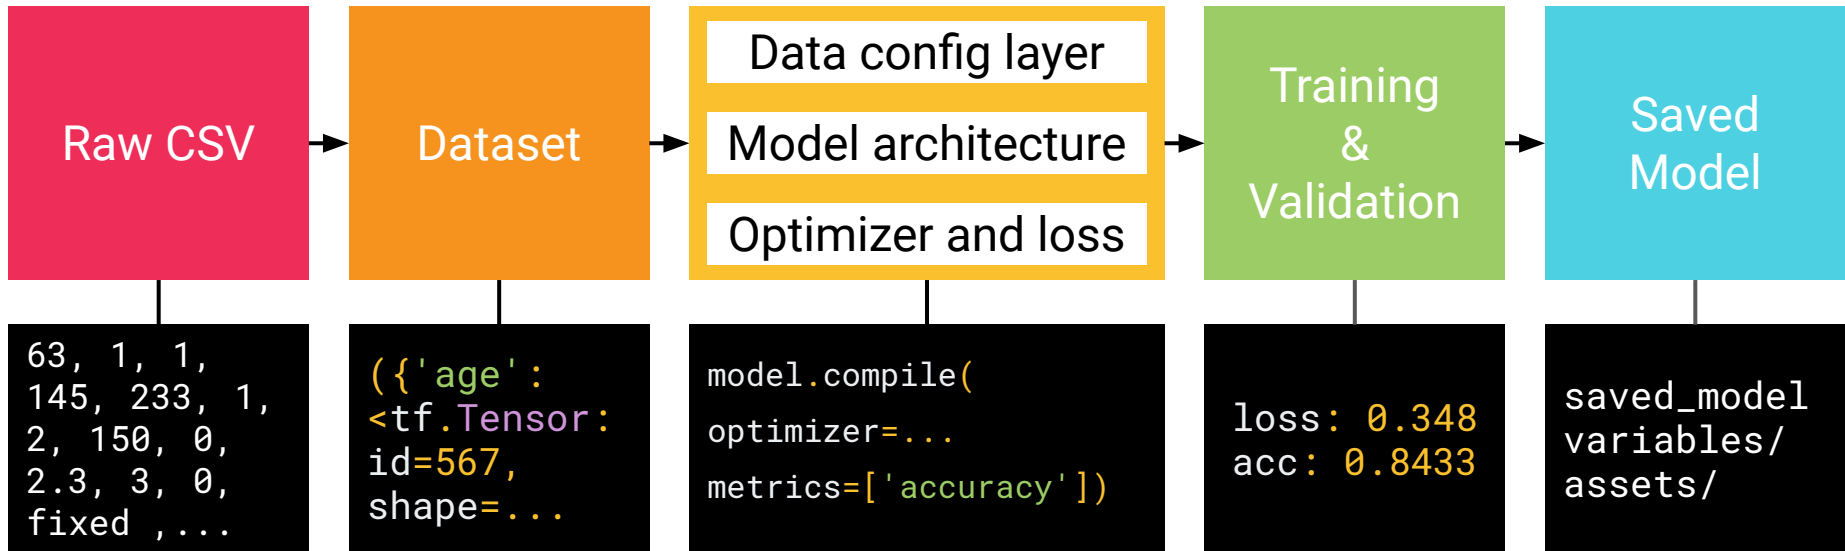


Saved
Model

Multi-stage Process



Multi-stage Process




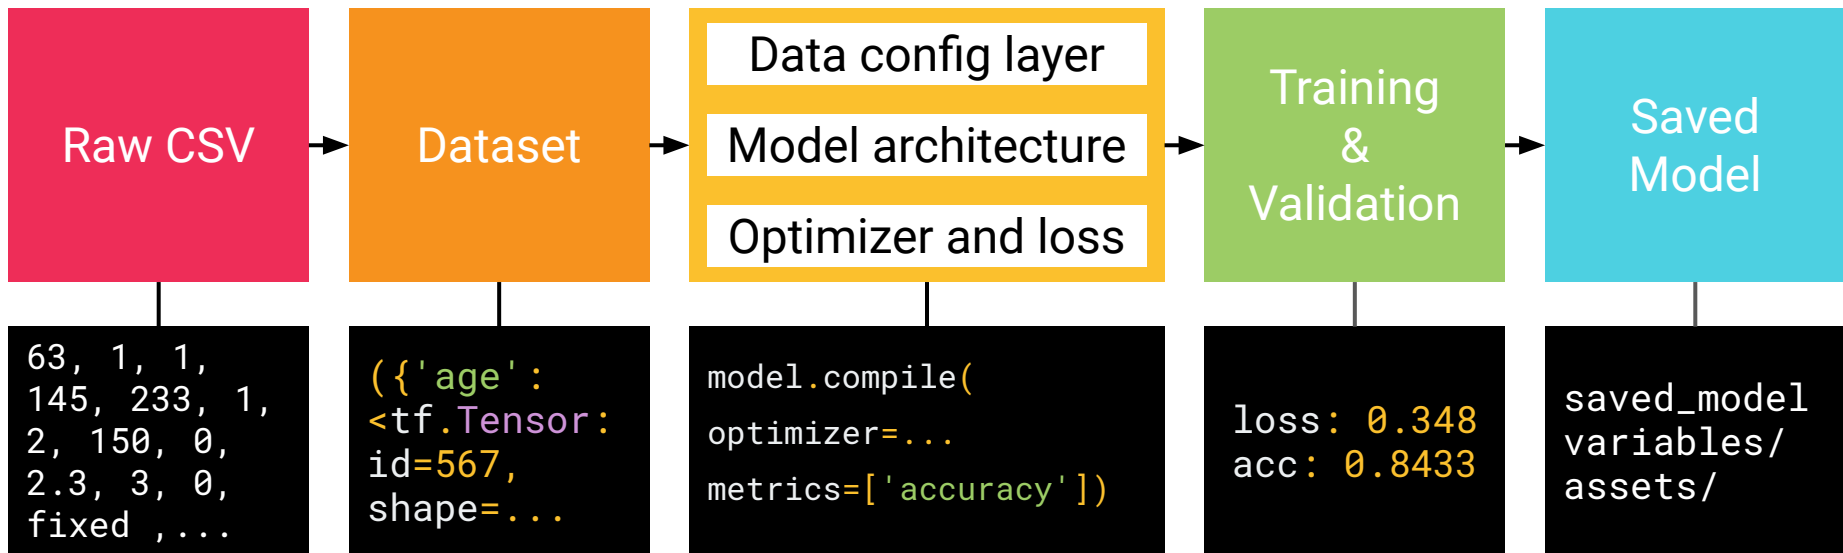


Conclusions

- Prototype with **Eager**.
- Preprocess with **Datasets**.
- Transform with **Feature Columns**.
- Build with **Keras**.
- Package with **SavedModel**.



Multi-stage Process



Exercise 4

Structured data



Exercise 4

Goals

- Start from raw data
- Create a model

Visit

bit.ly/tf-ws4a

Facets pair-code.github.io/facets/



Beyond Hello World



A few of my favorites

- Machine Translation
- Image Captioning (the decoder is similar!)
- DCGan and Pix2Pix



The docs are code

Tutorials on tf.org/alpha are

- Backed by a Jupyter Notebook
- Can be run directly in Colab

They automatically

- Install the right TensorFlow version
- Download a dataset
- Train a model
- Show you the result

tensorflow.org/alpha/tutorials/text/image_captioning

TensorFlow > Learn > TensorFlow Core > TF 2.0 Alpha

Image Captioning with Attention

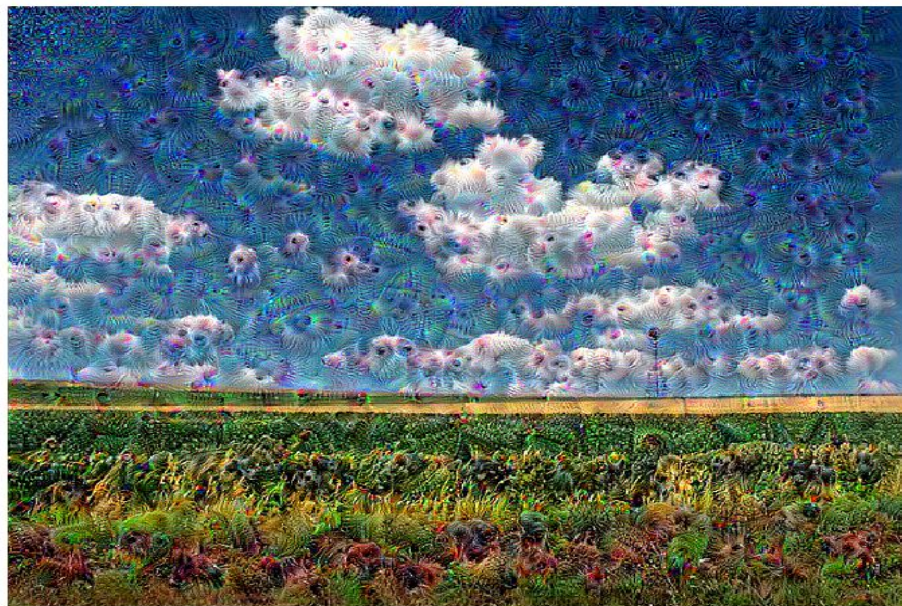
 Run in Google Colab



View source on GitHub


Given an image like the below, our goal is to generate a caption, such as "a surfer riding on a wave".





<https://github.com/random-forests/applied-dl/blob/master/examples/9-deep-dream-minimal.ipynb>

Code walkthrough



<https://github.com/random-forests/applied-dl/blob/master/examples/9-image-colorization.ipynb>

Exercise 6

Deep Dream



Is anyone bilingual? Trilingual?

When translating, do you...


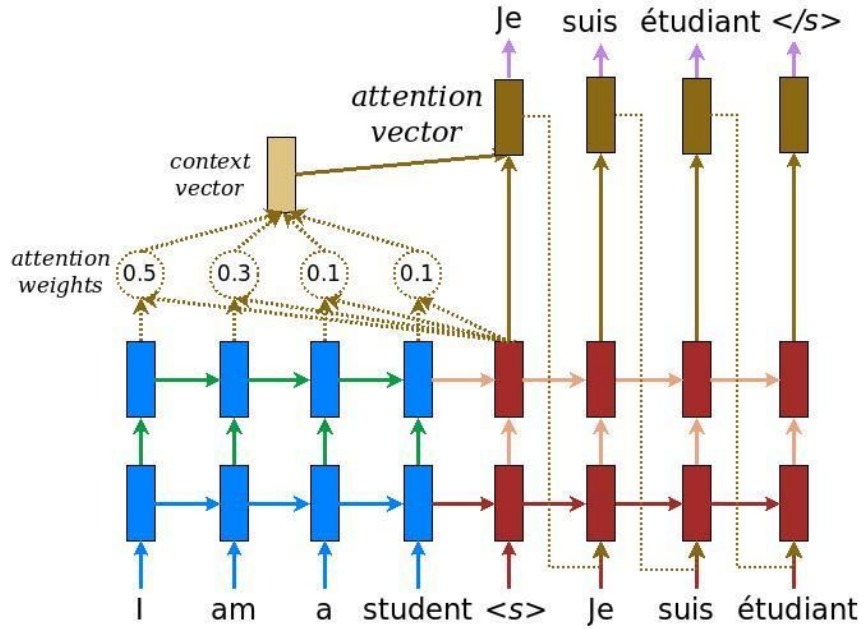
- Go directly from source -> target
- Or, go from source -> **intermediate representation** -> target.



Machine translation tutorials

- [Hello world](#) (seq2seq), trains in about a minute.
- [Neural Machine Translation with Attention](#)
- [Transformer](#)

P.S., isn't 2019 cool? It's **amazing** this is possible.



https://www.tensorflow.org/alpha/tutorials/sequences/nmt_with_attention



GENERATOR
"The Artist"
A neural network trying to
create pictures of cats that
look real.



DISCRIMINATOR
"The Art Critic"
A neural network examining
cat pictures to determine if
they're real or fake.




Thousands of real-world
images labeled "CAT"



<https://www.tensorflow.org/alpha/tutorials/generative/dcgan>


Input Image



Ground Truth



Predicted Image



<https://www.tensorflow.org/alpha/tutorials/generative/pix2pix>



Prediction Caption: the person is riding a surfboard in the ocean <end>

https://www.tensorflow.org/alpha/tutorials/sequences/image_captioning

Exercise 7

Seq2Seq



Exercise 6

Goals

- Use a pretrained CNN
- Extract intermediate activations
- Compute gradients w.r.t. an image

Visit

bit.ly/tf-ws6



Exercise 7

Goals

- Train an English to Spanish model, just for fun
- Learn about encoder / decoders

Visit

bit.ly/minimal-nmt



Under the hood

Let's make this faster

```
lstm_cell = tf.keras.layers.LSTMCell(10)
```

```
def fn(input, state):  
    return lstm_cell(input, state)
```

```
input = tf.zeros([10, 10]); state = [tf.zeros([10, 10])] * 2  
lstm_cell(input, state); fn(input, state) # warm up
```

```
# benchmark
```

```
timeit.timeit(lambda: lstm_cell(input, state), number=10) # 0.03
```

Let's make this faster

```
lstm_cell = tf.keras.layers.LSTMCell(10)
```

```
@tf.function
```

```
def fn(input, state):
```

```
    return lstm_cell(input, state)
```

```
input = tf.zeros([10, 10]); state = [tf.zeros([10, 10])] * 2
```

```
lstm_cell(input, state); fn(input, state) # warm up
```

```
# benchmark
```

```
timeit.timeit(lambda: lstm_cell(input, state), number=10) # 0.03
```

```
timeit.timeit(lambda: fn(input, state), number=10) # 0.004
```

AutoGraph makes this possible

```
@tf.function
```

```
def f(x):
```

```
    while tf.reduce_sum(x) > 1:
```

```
        x = tf.tanh(x)
```

```
    return x
```

```
# you never need to run this (unless curious)
```

```
print(tf.autograph.to_code(f))
```

Generated code

```
def tf__f(x):  
    def loop_test(x_1):  
        with ag__.function_scope('loop_test'):  
            return ag__.gt(tf.reduce_sum(x_1), 1)  
    def loop_body(x_1):  
        with ag__.function_scope('loop_body'):  
            with ag__.utils.control_dependency_on_returns(tf.print(x_1)):  
                tf_1, x = ag__.utils.alias_tensors(tf, x_1)  
                x = tf_1.tanh(x)  
            return x,  
    x = ag__.while_stmt(loop_test, loop_body, (x,), (tf,))  
    return x
```

Going big: tf.distribute.Strategy

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(64, input_shape=[10]),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')])

model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```


Going big: Multi-GPU

```
strategy = tf.distribute.MirroredStrategy()

with strategy.scope():
    model = tf.keras.models.Sequential([
        tf.keras.layers.Dense(64, input_shape=[10]),
        tf.keras.layers.Dense(64, activation='relu'),
        tf.keras.layers.Dense(10, activation='softmax')])

    model.compile(optimizer='adam',
                  loss='categorical_crossentropy',
                  metrics=['accuracy'])
```



What's different between TF1 and TF2?

Removed


- `session.run`
- `tf.control_dependencies`
- `tf.global_variables_initializer`
- `tf.cond`, `tf.while_loop`

Added

- `tf.function`, `AutoGraph`



TensorFlow.js




Demo #1


PoseNet



PoseNet

bit.ly/pose-net





Demo #2

BodyPix



BodyPix

bit.ly/body-pix





Learning more



Learn more

Tutorials and guides

- [tensorflow.org/alpha](https://www.tensorflow.org/alpha)

Books

- [Deep Learning with Python](#)
- Hands-On Machine Learning with Scikit-Learn and TensorFlow (version 2.0 is almost ready)

Courses

- [Intro to Deep Learning](#) (MIT)
- [Convolutional Neural Networks for Visual Recognition](#) (Stanford)



tf.thanks!

Josh Gordon (twitter.com/random_forests)



Extras



Reminders

- Deep learning as compression
- Interlingual representations



What's Deep Learning?

Representation learning

Automatic feature engineering

DATA


Which dataset do you want to use?




Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?

X_1



X_2



X_1^2



X_2^2



$X_1 X_2$



$\sin(X_1)$



$\sin(X_2)$



+


-

0 HIDDEN LAYERS


OUTPUT

Test loss 0.455

Training loss 0.455



Colors shows data, neuron and weight values.



☐ Show test data

☐ Discretize output

DATA


Which dataset do you want to use?




Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?

X_1



X_2



X_1^2



X_2^2



$X_1 X_2$



$\sin(X_1)$



$\sin(X_2)$



+


-

0 HIDDEN LAYERS

OUTPUT

Test loss 0.672

Training loss 0.627




Colors shows data, neuron and weight values.



☐ Show test data

☒ Discretize output





Feature engineering. What if we add a new feature:

$$z = x^2 + y^2.$$

Intuition

- All values of z will be positive.
- Yellow points are closer to the origin...
- So sum of their squared coords will be lower than red!



DATA


Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?



+


-

0 HIDDEN LAYERS


OUTPUT

Test loss 0.486

Training loss 0.543



Colors shows data, neuron and weight values.



☐ Show test data

☒ Discretize output

DATA

Which dataset do you want to use?



Ratio of training to test data: 50%



Noise: 0



Batch size: 10



REGENERATE

FEATURES

Which properties do you want to feed in?



+

-

1 HIDDEN LAYER

+

-

3 neurons




This is the output from one **neuron**. Hover to see it larger.

OUTPUT

Test loss 0.547

Training loss 0.600



Colors shows data, neuron and weight values.



☐ Show test data

☒ Discretize output

DATA


Which dataset do you want to use?




Ratio of training to test data: 50%



Noise: 0



Batch size: 10



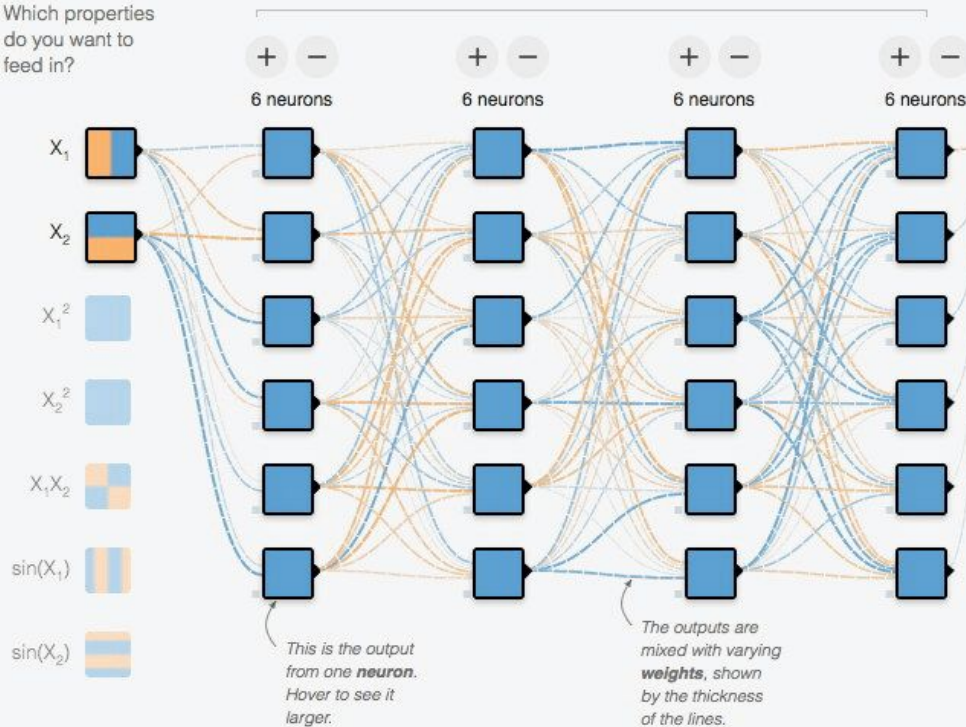
REGENERATE

FEATURES

Which properties do you want to feed in?




+ - 4 HIDDEN LAYERS




OUTPUT

Test loss 0.511

Training loss 0.502



Colors shows data, neuron and weight values.



☐ Show test data

☒ Discretize output

